

16-720 Computer Vision: Homework 5 (Spring 2023)

Neural Networks for Recognition

Instructor: Deva Ramanan

TAs: : Kangle Deng, Vidhi Jain, Xiaofeng Guo, Chung Hee Kim, Ingrid Navarro

Released on: April 12, 2023

Due: April 28, 2023

Instructions

1. **Integrity and collaboration:** Students are encouraged to work in groups but each student must submit their own work. Include the names of your collaborators in your write-up. Code should **NOT** be shared or copied. Please **DO NOT** use external code unless permitted. Plagiarism is prohibited and may lead to failure in this course.
2. Answer each question (for points) marked with a **Q** in the corresponding titled boxes.
3. **Start early!** This homework may take a long time to complete.
4. **Attempt to verify your implementation as you proceed.** If you don't verify that your implementation is correct on toy examples, you will risk having a huge mess when you put everything together.
5. **In your PDF, start a new page for each question and indicate the answer/page(s) correspondence carefully when submitting on Gradescope.** For some questions, this may leave a lot of blank space. If you skip a written question, just submit a blank page for it. This makes your work much easier to grade.
6. **Some questions will ask you to “Include your code in the writeup”.** For those questions, you can either copy/paste the code into a `verbatim` environment or include screenshots of your code.
7. If you have any questions or need clarifications, please post in Slack or visit the TAs during office hours.
8. **Submission:** The submission is on Gradescope, **you will be submitting both your writeup and code zip file.** The zip file, <`andrew-id.zip`> contains your code and any results files we ask you to save. **Note: You have to submit your write-up separately to Gradescope, and include results (and code snippets, when requested) in the write-up.**
9. **Do not** submit anything from the `data/` folder in your submission.
10. For your code submission, **do not** use any libraries other than `numpy`, `scipy`, `scikit-image`, `matplotlib` and (in the appropriate section) `pytorch`. Including other libraries (for example, `cv2`, etc.) **may lead to loss of credit** on the assignment.

11. To get the data for Section 3 onwards, we have included some scripts in `scripts`. For those who are on operating systems like Windows and cannot run the .sh script, you can also manually download by clicking on the link and unzipping the data. Download and unzip

http://www.cs.cmu.edu/~lkeselma/16720a_data/data.zip

http://www.cs.cmu.edu/~lkeselma/16720a_data/images.zip

and extract them to `data` and `image` folders.

12. For each coding question, refer to the comments inside the given Python scripts to see which function to implement. Insert your code into the designated place (where it says `your code here`), and **DO NOT** remove any of the given code elsewhere.
13. Feel free to use

`\codesection{name-of-file.py}`

to include the code for each question. Mark the pages for each question appropriately.

14. Increase the size of the answer blocks as required.
15. To include figures in the answer blocks, use only ‘`includegraphics`’ but not inside of a ‘`figure`’ block. This is because ‘`figure`’ is a floating environment, but you want to place your graphics in a particular place. Or feel free to define the `figure` block outside the `your_solution` block.

Contents

1 Theory	4
2 Implement a Fully Connected Network	8
2.1 Network Initialization	8
2.2 Forward Propagation	10
2.3 Backwards Propagation	12
2.4 Training Loop: Stochastic Gradient Descent	13
2.5 Numerical Gradient Checker	16
3 Training Models	20
4 Extract Text from Images	27
5 (Extra Credit) Image Compression with Autoencoders	38
5.1 Building the Autoencoder	38
5.2 Training the Autoencoder	42
5.3 Evaluating the Autoencoder	43
6 PyTorch	45
6.1 Train a neural network in PyTorch	46
6.2 Fine Tuning	63
6.3 Neural Networks in the Real World	70
Deliverables	76
References	77
Appendix	78
Mathematical overview	78
Backprop	79
Debugging Checklist	79

1 Theory

Q1.1 Theory [3 points] Prove that softmax is invariant to translation, that is

$$\text{softmax}(x) = \text{softmax}(x + c) \quad \forall c \in \mathbb{R}.$$

Softmax is defined as below, for each index i in a vector \mathbf{x} .

$$\text{softmax}(x_i) = \frac{e^{x_i}}{\sum_j e^{x_j}}$$

Often we use $c = -\max x_i$. Why is that a good idea? (Tip: consider the range of values that the numerator will have with $c = 0$ and $c = -\max x_i$)

Q1.1

1. Prove: As $\text{softmax}(x_i) = \frac{e^{x_i}}{\sum_j e^{x_j}}$, then $\text{softmax}(x_i + c) = \frac{e^{x_i+c}}{\sum_j e^{x_j+c}} = \frac{e^{x_i}*e^c}{\sum_j e^{x_j}*e^c}$

For $e^c \neq 0$, then divide both numerators and denominators by e^c , $\text{softmax}(x_i + c) = \frac{e^{x_i}}{\sum_j e^{x_j}} = \text{softmax}(x_i)$, so $\text{softmax}(x) = \text{softmax}(x + c) \quad \forall c \in \mathbb{R}$.

2. When $c=0$, e^{x_i} is from 0 to infinity. When $c = -\max x_i$, $e^{x_i-\max x_i}$ is from 0 to $e^0 = 1$. Then range is much smaller, which can help normalization and prevent the number from becoming too large.

Q1.2 Theory [3 points] Softmax can be written as a three-step process, with $s_i = e^{x_i}$, $S = \sum s_i$ and $\text{softmax}(x_i) = \frac{1}{S}s_i$.

1. As $x \in \mathbb{R}^d$, what are the properties of $\text{softmax}(x)$, namely what is the range of each element of the $\text{softmax}(x)$? What is the sum of all elements in the $\text{softmax}(x)$?

Q1.2.1

1. For $\text{softmax}(x_i) = \frac{1}{S}s_i = \frac{e^{x_i}}{\sum_i e^{x_i}}$, then range of each is from 0 to 1.

2. sum of all elements: $\sum_i \frac{e^{x_i}}{\sum_i e^{x_i}} = \frac{\sum_i e^{x_i}}{\sum_i e^{x_i}} = 1$

2. One could say that “softmax takes an arbitrary real valued vector x and turns it into a _____”.

Q1.2.2

From [Softmax function in Wikipedia](#)

softmax takes an arbitrary real valued vector x and turns it into a probability distribution of it.

3. What is the role of each step in the three-step process to compute $\text{softmax}(x)$? Explain.

Q1.2.3

1. get $s = e^{x_i} - c$: make the value in a fixed range. 2. get $S = \sum s_i$: get sum of values, so that it can be in probability distribution. 3. get each as $\frac{s}{S}$, make it in range from 0 to 1, as probability distribution.

Q1.3 Theory [3 points] Show that multi-layer neural networks without a non-linear activation function are equivalent to linear regression.

Q1.3

multi-layer neural networks: $y = f_n(w_n * x_n + b_n), x_n = f_{n-1}(w_{n-1} * x_{n-1} + b_{n-1}) \dots$
Without without a non-linear activation function: $y = k_n * (w_n * x_n + b_n) + g_n$ (k, g are constant)

$$y = k_n * w_n * x_n + k_n * b_n + g_n$$

this can be written as: $y = w'_n * x_n + b'_n$

Then for, $x_n = w'_{n-1} * x_{n-1} + b'_{n-1}$, so $y = w'_n * (w'_{n-1} * x_{n-1} + b'_{n-1}) + b'_n = w'_n * w'_{n-1} * x_{n-1} + w'_n * b'_{n-1} + b'_n$

This can be written as: $y = w' * x_{n-1} + b'$ Same for $n-2, \dots, 2, 1$.

Finally $y = w * x + b$, which is the same as linear regression

Q1.4 Theory [3 points] Given the sigmoid activation function $\sigma(x) = \frac{1}{1+e^{-x}}$, derive the gradient of the sigmoid function and show that it can be written as a function of $\sigma(x)$ (without having access to x directly)

Q1.4

$$\text{gradient: } \frac{d\sigma(x)}{dx} = \frac{1}{1+e^{-x}} = \frac{-(-e^{-x})}{(1+e^{-x})^2} = \frac{e^{-x}}{(1+e^{-x})^2}$$

In function of $\sigma(x)$:

$$\frac{d\sigma(x)}{dx} = \frac{e^{-x}}{1+e^{-x}} * \frac{1}{1+e^{-x}}$$

Since $\frac{1}{1+e^{-x}} = \sigma(x)$

$$\frac{e^{-x}}{1+e^{-x}} = \frac{e^{-x}+1-1}{1+e^{-x}} = 1 - \frac{1}{1+e^{-x}} = 1 - \sigma(x)$$

$$\text{So } \frac{d\sigma(x)}{dx} = \sigma(x) * (1 - \sigma(x))$$

Q1.5 Theory [12 points] Given $y = Wx + b$ (or $y_i = \sum_{j=1}^d x_j W_{ij} + b_i$), and the gradient of some loss J with respect y , show how to get $\frac{\partial J}{\partial W}$, $\frac{\partial J}{\partial x}$ and $\frac{\partial J}{\partial b}$. Be sure to do the derivatives with scalars and re-form the matrix form afterward. Here is some notional suggestions.

$$\frac{\partial J}{\partial y} = \delta \in \mathbb{R}^{k \times 1} \quad W \in \mathbb{R}^{k \times d} \quad x \in \mathbb{R}^{d \times 1} \quad b \in \mathbb{R}^{k \times 1}$$

We won't grade the derivative with respect to b but you should do it anyways, you will need it later in the assignment.

Q1.5

The shape of x, w, b 's gradient should be the same as shape of x, w, b

Also, apply chain rule here:

$$1. \frac{\partial J}{\partial W} = \frac{\partial J}{\partial y} * \frac{\partial y}{\partial W} = \delta * \frac{\partial y}{\partial W}$$

$$\text{Since } y = Wx + b, \text{ then } \frac{\partial(Wx+b)}{\partial W} = x$$

Because shape of w 's gradient should be the same as shape of w

$$\text{Then } \frac{\partial J}{\partial W} = \delta * x^T$$

$$2. \frac{\partial J}{\partial x} = \frac{\partial J}{\partial y} * \frac{\partial y}{\partial x} = \delta * \frac{\partial y}{\partial x}$$

$$\text{Since } y = Wx + b, \text{ then } \frac{\partial(Wx+b)}{\partial x} = W$$

Because shape of x 's gradient should be the same as shape of x ,

$$\text{Then } \frac{\partial J}{\partial x} = W^T * \delta$$

$$3. \frac{\partial J}{\partial b} = \frac{\partial J}{\partial y} * \frac{\partial y}{\partial b} = \delta * \frac{\partial y}{\partial b}$$

$$\text{Since } y = Wx + b, \text{ then } \frac{\partial(Wx+b)}{\partial b} = 1$$

Because shape of x 's gradient should be the same as shape of x ,


$$\text{Then } \frac{\partial J}{\partial b} = \delta$$

Q1.6 Theory [4 points] When the neural network applies the elementwise activation function (such as sigmoid), the gradient of the activation function scales the backpropagation update. This is directly from the chain rule, $\frac{d}{dx}f(g(x)) = f'(g(x))g'(x)$.

1. Consider the sigmoid activation function for deep neural networks. Why might it lead to a “vanishing gradient” problem if it is used for many layers (consider plotting Q1.4)?

Q1.6.1

Plot of sigmid and derivative of the Sigmoid from [Derivative of the Sigmoid function](#)



From the plot, we can see that the gradient of Sigmoid function is between 0 and 0.25, also with x increasing, the gradient will decrease and become near to 0. So with many layers, the value will become smaller and smaller until 0, which is also the vanishing gradient.

- Often it is replaced with $\tanh(x) = \frac{1-e^{-2x}}{1+e^{-2x}}$. What are the output ranges of both \tanh and sigmoid? Why might we prefer \tanh ?


Q1.6.2

- For sigmoid, the for e^{-x} is from 0 to +infinity, so $\frac{1}{1+e^{-x}}$ is from 0 to 1.
For \tanh , the range is from -1 to 1.
- We prefer \tanh for it has higher gradient value and wilder range.

- Why does $\tanh(x)$ have less of a vanishing gradient problem? (Hint: plotting the derivatives helps! For reference: $\tanh'(x) = 1 - \tanh(x)^2$)

Q1.6.3

Plot of $\tanh(x)$ and derivative of the $\tanh(x)$ from [Calculating the derivative of Tanh:](#)



We can see that the range is from 0 to 1, which is much larger than 0.25 with same x value. So it has less of a vanishing gradient problem

- \tanh is a scaled and shifted version of the sigmoid. Show how $\tanh(x)$ can be written in terms of $\sigma(x)$.

Q1.6.4

$$\tanh(x) = \frac{1-e^{-2x}}{1+e^{-2x}} = -\frac{e^{-2x}-1}{1+e^{-2x}} = -\frac{e^{-2x}+1-2}{1+e^{-2x}} = -\frac{e^{-2x}+1}{1+e^{-2x}} + 2 * \frac{1}{1+e^{-2x}} = -1 + 2 * \frac{1}{1+e^{-2x}}$$

Since $\sigma(x) = \frac{1}{1+e^{-x}}$
So, $\tanh(x) = -1 + 2 * \sigma(2x)$

2 Implement a Fully Connected Network

When implementing the following functions, make sure you run `python/run_q2.py` along the way to check if your implemented functions work as expected.

2.1 Network Initialization

Q2.1.1 Theory [3 points] Why is it not a good idea to initialize a network with all zeros? If you imagine that every layer has weights and biases, what can a zero-initialized network output be after training?

Q2.1.1

Reason: learned from slide and [How to Initialize Weights in Neural Networks?](#):
If every layer has same weights and biases. Every neuron will do the same work and learn the same features of their inputs.
Then there will be no difference for each one. Then the result can be very bad.
Besides, giving different initial value can prevent situation that all neurons are in local minima.

Q2.1.2 Code [3 points] In `python/nn.py`, implement a function to initialize one layer's weights with Xavier initialization [1], where $Var[w] = \frac{2}{n_{in}+n_{out}}$ where n is the dimensionality of the vectors and you use a **uniform distribution** to sample random numbers (see eq 16 in [Glorot et al]). **Include your code in the write-up.**

Q2.1.2

```
1 def initialize_weights(in_size,out_size,params,name=''):
2     W, b = None, None
3     ##### your code here #####
4     #####
5     #####
6     # W = np.zeros((in_size,out_size))
7     b = np.zeros((out_size))
8     W_size = (in_size,out_size)
9     # according to formula in paper:
10    high = np.sqrt(6)/np.sqrt(in_size+out_size)
11    low = - high
12    W = np.random.uniform(low,high,(W_size))
13
14    params['W' + name] = W
15    params['b' + name] = b
```

Q2.1.3 Theory [2 points] Why do we initialize with random numbers? Why do we scale the initialization depending on layer size (see Fig 6 in the [Glorot et al])?

Q2.1.3

1. For it can prevent that every neuron from doing the same work and learning the same features of their inputs.
2. From Fig 6 we can see that when initialization depending on layer size, the activation value in different will be similar. This can help delete the difference of different layers

2.2 Forward Propagation

The appendix (sec 6.3) has the math for forward propagation, we will implement it here.

Q2.2.1 Code [4 points] In python/numpy.py, implement sigmoid, along with forward propagation for a single layer with an activation function, namely $y = \sigma(XW + b)$, returning the output and intermediate results for an $N \times D$ dimension input X , with examples along the rows, and data dimensions along the columns. **Include your code in the write-up.**

Q2.2.1

```
1 def sigmoid(x):
2     res = None
3     ##### your code here #####
4     #####
5     res = 1/(1+np.exp(-x))
6     return res
7
8
9 def forward(X,params,name='',activation=sigmoid):
10    pre_act, post_act = None, None
11    # get the layer parameters
12    W = params['W' + name]
13    b = params['b' + name]
14    #####
15    ##### your code here #####
16    #####
17    # print(X.shape,W.shape,b.shape)
18    num = X.shape[0]
19    out_size = b.shape[0]
20    pre_act = np.zeros((num, out_size))
21    for i in range(num):
22        value = X[i] @ W + b
23        pre_act[i, :] = value
24
25    post_act = activation(pre_act)
26
27    # store the pre-activation and post-activation values
28    # these will be important in backprop
29    params['cache_' + name] = (X, pre_act, post_act)
30
31    return post_act
```

Q2.2.2 Code [3 points] In `python/nn.py`, implement the softmax function. Be sure to use the numerical stability trick you derived in Q1.1 softmax. **Include your code in the write-up.**

```
Q2.2.2

1 def softmax(x):
2     res = None
3
4     ##### your code here #####
5     #####
6
7     # print(x.shape)
8     N, n = x.shape
9     res = np.zeros((N, n))
10    for i in range(N):
11        # for each row
12        # use the numerical stability trick in Q1.1 softmax
13        c = - max(x[i, :])
14        sum = 0
15        for j in range(n):
16            res[i, j] = np.exp(x[i, j] + c)
17            sum += res[i, j]
18        res[i, :] = res[i, :] / sum
19    # print(res)
20
21    return res
```

Q2.2.3 Code [3 points] In `python/nn.py`, write a function to compute the accuracy of a set of labels, along with the scalar loss across the data. The loss function generally used for classification is the cross-entropy loss.

$$L_f(\mathbf{D}) = - \sum_{(\mathbf{x}, \mathbf{y}) \in \mathbf{D}} \mathbf{y} \cdot \log(f(\mathbf{x}))$$

Here \mathbf{D} is the full training dataset of data samples \mathbf{x} ($N \times 1$ vectors, N = dimensionality of data) and labels \mathbf{y} ($C \times 1$ one-hot vectors, C = number of classes), and $\mathbf{f} : \mathbb{R}^N \rightarrow [0, 1]^C$ is the classifier. The log is the natural log. **Include your code in the write-up.**

Q2.2.3

```
1 def compute_loss_and_acc(y, probs):
2     loss, acc = None, None
3
4     ##### your code here #####
5     #####
6
7     # print(y)
8     loss_matrix = y * np.log(probs)
9     loss = -np.sum(loss_matrix)
10    N = y.shape[0]
11    correct_num = 0
12    for i in range(N):
13        prob_index = np.argmax(probs[i, :])
14        y_index = np.argmax(y[i, :])
15        if y_index == prob_index:
16            correct_num += 1
17    acc = correct_num / N
18
19    return loss, acc
```

2.3 Backwards Propagation

Q2.3 Code [7 points] In `python/nn.py`, write a function to compute backpropagation for a single layer, given the original weights, the appropriate intermediate results, and the given gradient with respect to the loss. You should return the gradient with respect to X so you can feed it into the next layer. As a size check, your gradients should be the same

dimensions as the original objects. **Include your code in the write-up.**

Q2.3

```
1 def backwards(delta, params, name='', activation_deriv=sigmoid_deriv):
2     grad_X, grad_W, grad_b = None, None, None
3     # everything you may need for this layer
4     W = params['W' + name]
5     b = params['b' + name]
6     X, pre_act, post_act = params['cache_' + name]
7
8     # do the derivative through activation first
9     # (don't forget activation_deriv is a function of post_act)
10    # then compute the derivative W, b, and X
11    #####
12    ##### your code here #####
13    #####
14    # do the derivative through activation
15    action_deriv = activation_deriv(post_act)
16    # derivative of cross-entropy:delta
17    # print(action_deriv.shape)
18    # print(delta.shape)
19    # print(W.shape)
20    # print(b.shape)
21    # print(X.shape)
22    deriv_f = action_deriv * delta
23    #  $d(f(wx+b))/dw = x*f'$ 
24    grad_W = X.T @ deriv_f
25    #  $d(f(wx+b))/db = I*f'$ 
26    N = deriv_f.shape[0]
27    I = np.ones((1, N))
28    grad_b = (I @ deriv_f).reshape(b.shape)
29    #  $d(f(wx+b))/dx = w*f'$ 
30    grad_X = (W @ deriv_f.T).T
31
32    # store the gradients
33    params['grad_W' + name] = grad_W
34    params['grad_b' + name] = grad_b
35    return grad_X
```

2.4 Training Loop: Stochastic Gradient Descent

Q2.4 Code [5 points] In `python/nn.py`, write a function that takes the entire dataset (`x` and `y`) as input and splits it into random batches. In `python/run_q2.py`, write a training

loop that iterates over the batches, does forward and backward propagation, and applies a gradient update. The provided code samples batch only once, but it is also common to sample new random batches at each epoch. You may optionally try both strategies and note any difference in performance. **Include your code in the write-up.**

Q2.4(part1)

```
1 def get_random_batches(x,y,batch_size):
2     batches = []
3     ##### your code here #####
4     #####
5     # print(x.shape,y.shape)
6     N = x.shape[0]
7     num_per_batch = N // batch_size
8     rand_index = np.random.permutation(N)
9     start_index = 0
10    # print(x)
11    # print(rand_index)
12    # the first n-1 batches
13    for i in range(batch_size - 1):
14        batch_index = rand_index[start_index:start_index + num_per_batch]
15        start_index += num_per_batch
16        batch_x = x[batch_index]
17        batch_y = y[batch_index]
18        batches.append((batch_x, batch_y))
19    # the last batch(the rest)
20    batch_index = rand_index[start_index:]
21    batch_x = x[batch_index]
22    batch_y = y[batch_index]
23    # print(batch_x)
24    batches.append((batch_x, batch_y))
25
26    return batches
```

code for training loop:

```
1 for itr in range(max_iters):
2     total_loss = 0
3     avg_acc = 0
4     for xb,yb in batches:
5         ##### your code here #####
6         #####
7         #pass
8         # forward
9         h1 = forward(xb, params, 'layer1')
10        probs = forward(h1, params, 'output', softmax)
11        # loss
12        # be sure to add loss and accuracy to epoch totals
13        loss, acc = compute_loss_and_acc(yb, probs)
14        total_loss += loss
```

Q2.4(part2)

```
1     avg_acc += acc / batch_num
2
3     # backward
4     delta1 = probs - yb
5     # apply gradient
6     # gradients should be summed over batch samples
7     # Implement backwards!
8     delta2 = backwards(delta1, params, 'output', linear_deriv)
9     backwards(delta2, params, 'layer1', sigmoid_deriv)
10    # apply gradient
11    # for layers
12    for name in ('output', 'layer1'):
13        grad_W = params['grad_W' + name]
14        grad_b = params['grad_b' + name]
15        params['W' + name] -= learning_rate * grad_W
16        params['b' + name] -= learning_rate * grad_b
17
18    # gradients should be summed over batch samples
19
20    if itr % 100 == 0:
21        print("itr: {:02d} \t loss: {:.2f} \t acc :
22             {:.2f}".format(itr, total_loss, avg_acc))
```

2.5 Numerical Gradient Checker

Q2.5 Code [5 points] In `python/run_q2.py`, implement a numerical gradient checker. Instead of using the analytical gradients computed from the chain rule, add ϵ offset to each element in the weights, and compute the numerical gradient of the loss with central

differences. Central differences is just $\frac{f(x+\epsilon)-f(x-\epsilon)}{2\epsilon}$. Remember, this needs to be done for each scalar dimension in all of your weights independently. This should help you check your gradient code, so there is no need to show the result, but do **include your code in the writeup**.

Q2.5(part1)

```
1  eps = 1e-6
2  for k,v in params.items():
3      if '_' in k:
4          continue
5
6      ##### your code here #####
7      #####
8
9  if 'W' in k:
10     # for each value inside the parameter
11     n, m = v.shape
12     for i in range(n):
13         for j in range(m):
14             # add epsilon
15             value = params_orig[k][i, j]
16             value += eps
17             # run the network
18             # copy and replace
19             new_params = copy.deepcopy(params_orig)
20             new_params[k][i, j] = value
21             h1 = forward(x, new_params, 'layer1')
22             probs = forward(h1, new_params, 'output',
23                             ↪ softmax)
24             # get the loss
25             loss1, acc1 = compute_loss_and_acc(y, probs)
26             # subtract 2*epsilon
27             value -= eps * 2
28             new_params = copy.deepcopy(params_orig)
29             new_params[k][i, j] = value
30             h1 = forward(x, new_params, 'layer1')
31             probs = forward(h1, new_params, 'output',
32                             ↪ softmax)
33             loss2, acc2 = compute_loss_and_acc(y, probs)
34             # restore the original parameter value
35             # compute derivative with central diffs
36             cen_diff = (loss1 - loss2) / (2 * eps)
37             params['grad_' + k][i, j] = cen_diff
38
39     else:
40         n = v.shape[0]
41         for i in range(n):
42             # add epsilon
43             value = params_orig[k][i]
44             value += eps
45             # run the network
46
47 see the rest code in next page
```

Q2.5(part2)

```
1      # copy and replace
2      new_params = copy.deepcopy(params_orig)
3      new_params[k][i] = value
4      h1 = forward(x, new_params, 'layer1')
5      probs = forward(h1, new_params, 'output', softmax)
6      # get the loss
7      loss1, acc1 = compute_loss_and_acc(y, probs)
8      # subtract 2*epsilon
9      value -= eps * 2
10     new_params = copy.deepcopy(params_orig)
11     new_params[k][i] = value
12     h1 = forward(x, new_params, 'layer1')
13     probs = forward(h1, new_params, 'output', softmax)
14     loss2, acc2 = compute_loss_and_acc(y, probs)
15     # restore the original parameter value
16     # compute derivative with central diffs
17     cen_diff = (loss1 - loss2) / (2 * eps)
18     params['grad_' + k][i] = cen_diff
```

3 Training Models

Follow instructions to download the data in `data` and `image` folders.

Since our input images are 32×32 images, unrolled into one 1024-dimensional vector, that gets multiplied by $\mathbf{W}^{(1)}$, each row of $\mathbf{W}^{(1)}$ can be seen as a weight image. Reshaping each row into a 32×32 image can give us an idea of what types of images each unit in the hidden layer has a high response to.

We have provided you three data `.mat` files to use for this section.

The training data in `nist36_train.mat` contains samples for each of the 26 upper-case letters of the alphabet and the 10 digits. This is the set you should use for training your network.

The cross-validation set in `nist36_valid.mat` contains samples from each class, and should be used in the training loop to see how the network is performing on data that it is not training on. This will help to spot overfitting.

Finally, the test data in `nist36_test.mat` contains testing data, and should be used for the final evaluation of your best model to see how well it will generalize to new unseen data.

Use `python/run_q3.py` for this question, and refer to the comments for what to implement.

Q3.1 Code [5 points] Train a network from scratch. Use a single hidden layer with 64 hidden units, and train for at least 50 epochs. The script will generate two plots:

- (1) the accuracy on both the training and validation set over the epochs, and
- (2) the cross-entropy loss averaged over the data.

Tune the batch size and learning rate for accuracy on the validation set of at least 75%. Include the plots in your write-up. Hint: Use fixed random seeds to improve reproducibility.

Q3.1



result:(fixed random seed:np.random.seed(1))(batch_size = 100 learning_rate = 2e-3)

Validation accuracy: 0.7608333333333334

Test accuracy: 0.7605555555555555

plot:

1.acc in left and 2.loss in the right:



Code:

```
1  for xb,yb in batches:
2      # training loop can be exactly the same as q2!
3      #####
4      ##### your code here #####
5      #####
6      # img = np.transpose(xb[0].reshape((32,32)))
7      # plt.imshow(img)
8      # plt.show()
9      h1 = forward(xb, params, 'layer1')
10     probs = forward(h1, params, 'output', softmax)
11     # loss
12     # be sure to add loss and accuracy to epoch totals
13     loss, acc = compute_loss_and_acc(yb, probs)
14     total_loss += loss
15     avg_acc += acc / batch_num
16     # backward
17     # the derivative of cross-entropy(softmax(x)) is probs - 1[correct
18     # actions]
19     delta1 = probs - yb
20     delta2 = backwards(delta1, params, 'output', linear_deriv)
21     # Implement backwards!
22     backwards(delta2, params, 'layer1', sigmoid_deriv)
23     # apply gradient
24     # for layers
25     for name in ('output', 'layer1'):
26         grad_W = params['grad_W' + name]
27         grad_b = params['grad_b' + name]
28         params['W' + name] -= learning_rate * grad_W
29         params['b' + name] -= learning_rate * grad_b
```


Q3.2 Tune the Learning rate - Writeup [3 points] Use the script in Q3.1 to train and generate accuracy and loss plots for each of these three networks:

- (1) one with your tuned learning rate,
- (2) one with 10 times that learning rate, and
- (3) one with one-tenth that learning rate.


Include total of six plots in your write-up. Comment on how the learning rates affect the training, and report the final accuracy of the best network on the test set. *Hint: Use fixed random seeds to improve reproducibility.*

Q3.2


fixed random seed:`np.random.seed(1)`
 (1)tuned learning rate:`learning_rate = 2e-3`:
 Validation accuracy: 0.760833333333334
 Test accuracy: 0.7605555555555555
 plot:
 1.acc in left and 2.loss in the right:



(2) with 10 times that learning rate: 2e-2:
 1.acc in left and 2.loss in the right:



(3) one-tenth that learning rate: 2e-4:
 1.acc in left and 2.loss in the right:



Comment: for $w = w - lr * dw$, if learning rate is too large, the parameters will change sharply in each epoch, the plot is also sharp, then it may jump around the best value, which makes training bad. If it's too small, the change will be little, the plot can be smooth, then it may take lots of more steps to reach the best values and the training is also not that good.

Report:


Validation accuracy: 0.7705555555555555

Test accuracy: 0.7677777777777778 23


Q3.3 Learned Weights - Writeup [3 points] The script will visualize the first layer weights as 64 32x32 images, both immediately after initialization and after full training. Include both visualizations in your write-up. Comment on the learned weights and compare them to the initialized weights. Do you notice any patterns?

Q3.3

before training:



after training:




comment on learned weights and compare them to the initialized weights:
For the initialized weights are random, it looks like with lots of noise. After training
the weights look like very smooth. The values nearby are very close.
patterns: values become similar in closed area and smooth.

Q3.4 Confusion Matrix - Writeup [3 points] Visualize and include the confusion matrix of the test data for your best model. Comment on the top few pairs of classes that are most commonly confused.

Q3.4

confusion matrix:



The most commonly confused pairs: 'O' with '0', 'O' with 'D', 'Z' with '2'. Because they look very similar in shape. It's easy to be confused

4 Extract Text from Images




Figure 1: Sample image with handwritten characters annotated with boxes around each character.

Now that you have a network that can recognize handwritten letters with reasonable accuracy, you can now use it to parse text in an image. Given an image with some text on it, our goal is to have a function that returns the actual text in the image. However, since your neural network expects a binary image with a single character, you will need to process the input image to extract each character. There are various approaches that can be done so feel free to use any strategy you like.

Here we outline one possible method, another is given in a [tutorial](#)

1. Process the image ([blur](#), [threshold](#), [opening morphology](#), etc., perhaps in that order) to classify all pixels as being part of a character or background.
2. Find connected groups of character pixels (see [skimage.measure.label](#)). Place a bounding box around each connected component.
3. Group the letters based on which line of the text they are a part of, and sort each group so that the letters are in the order they appear on the page.
4. Take each bounding box one at a time and resize it to 32×32 , classify it with your network, and report the characters in order (inserting spaces when it makes sense).


Since the network you trained likely does not have perfect accuracy, you can expect there to be some errors in your final text parsing. Whichever method you choose to implement for character detection, you should be able to place a box on most of the characters in the image. We have provided you with `01_list.jpg`, `02_letters.jpg`, `03_haiku.jpg` and `04_deep.jpg` to test your implementation on.

Q4.1 Failures due to Assumptions - Theory [4 points] The method outlined above is pretty simplistic, and while it works for the given text samples, it makes several assumptions. What are two big assumptions that the sample method makes? In your writeup, include two example images where you expect the character detection to fail (for example, miss valid letters, misclassify letters or respond to non-letters).

Q4.1

Assumption1: The characters are in good shape and clear which means it's complete and every feature is kept. If it's written fast, some features may be lost or extra features are added:

Example the characters not in good shape and lose features or have extra wrong feature:



Assumption2 and 3: All letters have font forms similar to the training set letter font types and are all uppercase letters:

Examples that may be wrong(lowercase letters or rare font):

A B C D E F G H I J
K L M N O P Q R S
T U V W X Y Z
a b c d e f g h i j k l m n o p q r s t u v w x y z

Q4.2 Find Letters - Code [10 points] In `python/q4.py`, implement the function to find letters in the image. Given an RGB image, this function should return bounding boxes for all of the located handwritten characters in the image, as well as a binary black-and-white version of the image `im`. Each row of the matrix should contain `[y1, x1, y2, x2]` the positions of the top-left and bottom-right corners of the box. The black-and-white image should be between 0.0 to 1.0, **with the characters in black and the background in white**. Hint: Since we read text left to right, top to bottom, we can use this to cluster the coordinates. **Include your code in the write-up.**



Q4.2

```
1 def findLetters(image):
2     bboxes = []
3     bw = None
4     ##### your code here #####
5     #####
6     #####
7     # estimate noise
8     # denoise
9     denoised = skimage.restoration.denoise_bilateral(image,channel_axis=-1|,
10           ↳ sigma_color=0.05, sigma_spatial=15)
11     # -> greyscale
12     grayscale = skimage.color.rgb2gray(denoised)
13     # -> threshold
14     thresh = skimage.filters.threshold_otsu(grayscale)
15     binary = grayscale < thresh
16     # -> morphology
17     #tophat = skimage.morphology.black_tophat(binary)
18     dilation =skimage.morphology.binary_dilation(binary)
19     dilation = skimage.morphology.binary_dilation(dilation)
20     dilation = skimage.morphology.binary_dilation(dilation)
21     dilation = skimage.morphology.binary_dilation(dilation)
22     # -> label
23     label_image = skimage.measure.label(dilation)
24     regions = skimage.measure.regionprops(label_image)
25     # -> skip small boxes
26     # check
27     #skimage.io.imshow(denoised)
28     #skimage.io.imshow(binary)
29     # skimage.io.imshow(label_image)
30     # skimage.io.imshow(dilation)
31     areas=[]
32     for region in regions:
33         areas.append(region.area)
34     max_area = np.max(areas)
35     for region in regions:
36         #if region.area >= 500:
37         if region.area >= max_area*0.3:
38             bboxes.append(region.bbox)
39             # with the characters in black and the background in white.
40             # should be between 0.0 to 1.0
41             dilation = ~dilation
42             dilation=dilation.astype(float)
43             bw = dilation
44
45             return bboxes, bw
```



Q4.3 Eval Bounding Box Accuracy - Writeup [5 points] Using `python/run_q4.py`, visualize all of the located boxes on top of the binary image to show the accuracy of your `findLetters(..)` function. Include all the resulting images in your write-up.

Q4.3

Result1 and 2:



Result3 and 4:



We can see that all letters are recognized correctly.

Q4.4 Eval Letter Detection - Code/Writeup [8 points] In `python/run_q4.py`, you will now load the image, find the character locations, classify each one with the network you trained in **Q3.1**, and return the text contained in the image. Be sure you try to make your detected images look like the images from the training set. Visualize them and act accordingly. If you find that your classifier performs poorly, consider dilation under skimage morphology to make the letters thicker.

Your solution is correct if you can correctly detect approximately 100% of the letters and classify approximately 80% of the letters in each of the sample images.

Run your `run_q4` on all of the provided sample images in `images/`. Include the extracted text in your writeup. It is fine if your code ignores spaces, but if so, please add them manually

in the writeup.

Q4.4(part1)

Result:

```
T0 D0 LIST
I MA8E A T0 DO LIST
2 CH5EK OFF THE FIRST
3HING QN TO D0 LIST
3 R5ALIZE YOU HAVE ALREADYT
C0MPL5T5D 2 THINGS
4 REWARD YOURSELF WITB
A NAP

2 BC D E F G
H IJ K L M N
0 P Q R S T U
V W X Y Z
1 Z 3 F S 67 I9 0

HAIKUS ARE EASY
BUT SQMETIMES TREX DDNT MAKR SENGE
REFRIGERATOR

J55P LEARNING
D55P5F LEAFHING
J55FJST IEARNIHG
```

Image1:

T0 D0 LIST
I MA8E A T0 DO LIST
2 CH5EK OFF THE FIRST
3HING QN TO D0 LIST
3 R5ALIZE YOU HAVE ALREADYT
C0MPL5T5D 2 THINGS
4 REWARD Y0URSELF WITB
A NAP

Accuracy = 98/115 = 85.2%

Image2:

2 B C D E F G
H I J K L M N
0 P Q R S T U
V W X Y Z
1 Z 3 F S 6 7 I 9 0

Accuracy = 30/36 = 83.3%

Image3:

HAIKUS ARE EASY
BUT SQMETIMES TREX DDNT MAKR SENGE
REFRIGERATOR

Accuracy = 49/54 = 90.7%

(see the rest in next page)

Q4.4(part2)

Image4:

J55P LEARNING

D55P5F LEAFHING

J55FJST IEARNIHG

Accuracy = 27/41 = 65.9%

Overall accuracy = 204/246= 82.9% It's larger than 80%
code:

```
1 # find the rows using..RANSAC, counting, clustering, etc.
2 #####
3 ##### your code here #####
4 #####
5 # Group the letters based on which line of the text they are a part of, and sort
6 # each
7 # group so that the letters are in the order they appear on the page.
8 row_lists = []
9 # box: minr, minc, maxr, maxc
10 # sort with row
11 bboxes.sort(key=lambda x: x[0])
12 #h = bboxes[0][2] - bboxes[0][2]
13 threshold_row = 100
14 row_judge = bboxes[0][0]
15 #print(row_num)
16 #print(len(bboxes))
17 row_list = []
18 for bbox in bboxes:
19     row = bbox[0]
20     if abs(row - row_judge) < threshold_row:
21         row_list.append(bbox)
22     else:
23         # sort each group with col and append
24         row_list.sort(key=lambda x: x[1])
25         row_lists.append(row_list)
26         # start new row
27         row_list = []
28         row_list.append(bbox)
29         #update new
30         row_judge = row
31 # last row
32 row_list.sort(key=lambda x: x[1])
33 row_lists.append(row_list)
34 #print(row_lists)
#(see the rest code in next page)
```

Q4.4(part2)

```
1 # consider doing a square crop, and even using np.pad() to get your images
2     ↵ looking more like the dataset
3 #####
4 ##### your code here #####
5 #####
6 chars=[]
7 counts = []
8 for row_list in row_lists:
9     count = 0
10    for bbox in row_list:
11        count+=1
12        minr, minc, maxr, maxc = bbox
13        crop_img = bw[minr:maxr,minc:maxc]
14        # plt.imshow(crop_img)
15        # plt.show()
16        # square crop, and even using np.pad()
17        h = maxr-minr
18        w = maxc-minc
19        diff = int(abs(h-w)//2)
20        pad_h = int(h/7)
21        pad_w = int(h / 7)
22        if h > w:
23            pad_img = np.pad(crop_img, ((pad_h, pad_h), (diff+pad_w,
24                                         ↵ diff+pad_w)), 'maximum')
25        else:
26            pad_img = np.pad(crop_img, ((diff+pad_h,
27                                         ↵ diff+pad_h),(pad_w, pad_w)), 'maximum')
28        # plt.show() '
29        # resize: input images are 32 × 32 images
30        resized_img = skimage.transform.resize(pad_img, (32, 32))
31        # before you flatten, transpose the image
32        trans_img = np.transpose(resized_img)
33        #plt.imshow(resized_img)
34        #plt.show()
35        flatten_img = trans_img.flatten()
36        chars.append(flatten_img)
37        counts.append(count)
38 #print(counts)
39 see the rest in next part
```

Q4.4(part3)

```
1 ##### your code here #####
2 ##### your code here #####
3 ##### your code here #####
4 char_with = bboxes[0][3] - bboxes[0][1]
5 #print(char_with)
6 line_num = len(counts)
7 chars = np.array(chars)
8 h1 = forward(chars, params, 'layer1')
9 probs = forward(h1, params, 'output', softmax)
10 res = []
11 k=0
12 count = counts[k]
13 for prob in probs:
14     if counts[k]-count>0 and \
15         (row_lists[k][counts[k]-count][1] -
16          row_lists[k][counts[k]-count-1][3])>char_with*0.8:
17         res.append(' ')
18     index = np.argmax(prob)
19     letter = letters[index]
20     res.append(letter)
21     count -= 1
22     if count == 0:
23         res.append('\n')
24         k += 1
25         if k < line_num:
26             count = counts[k]
27 print(''.join(res))
```

5 (Extra Credit) Image Compression with Autoencoders

An autoencoder is a neural network that is trained to attempt to copy its input to its output, but it usually allows copying only approximately. This is typically achieved by restricting the number of hidden nodes inside the autoencoder; in other words, the autoencoder would be forced to learn to *represent* data with this limited number of hidden nodes. This is a useful way of learning compressed representations. In this section, we will continue using the NIST36 dataset you have from the previous questions. Use `python/run_q5.py` for this question.

5.1 Building the Autoencoder

Q5.1.1 (Extra Credit) Code [5 points] Due to the difficulty in training auto-encoders, we have to move to the $relu(x) = \max(x, 0)$ activation function. It is provided for you in `util.py`. Implement an autoencoder where the layers are

- 1024 to 32 dimensions, followed by a ReLU
- 32 to 32 dimensions, followed by a ReLU
- 32 to 32 dimensions, followed by a ReLU
- 32 to 1024 dimensions, followed by a sigmoid (this normalizes the image output for us)

The loss function that you're using is the total squared error for the output image compared to the input image (they should be the same!). **Include your code in the writeup.**

Q5.1.1

```
1 # initialize layers here
2 #####
3 ##### your code here #####
4 #####
5
6 initialize_weights(1024, 32, params, "layer1")
7 # zero-initialized momentum accumulators
8 params['Mw' + 'layer1'] = np.zeros((1024,32))
9 params['Mb' + 'layer1'] = np.zeros((32))
10
11 initialize_weights(32, 32, params, "layer2")
12 params['Mw' + 'layer2'] = np.zeros((32,32))
13 params['Mb' + 'layer2'] = np.zeros((32))
14
15 initialize_weights(32, 32, params, "layer3")
16 params['Mw' + 'layer3'] = np.zeros((32,32))
17 params['Mb' + 'layer3'] = np.zeros((32))
18
19 initialize_weights(32, 1024, params, "output")
20 params['Mw' + 'output'] = np.zeros((32,1024))
21 params['Mb' + 'output'] = np.zeros((1024))
22 # should look like your previous training loops
23
24 #####
25 ##### your code here #####
26 #####
27 # forward
28 h = forward(xb, params, 'layer1',relu)
29 h = forward(h, params, 'layer2',relu)
30 h = forward(h, params, 'layer3',relu)
31 probs = forward(h, params, 'output', sigmoid)
32 # squared error for the output image compared to the input image
33 loss = np.sum((xb-probs)**2)
34 total_loss+=loss
35 # backward
36 delta1 = -2*(xb-probs)
37 delta2 = backwards(delta1, params, 'output', sigmoid_deriv)
38 delta3 = backwards(delta2, params, 'layer3', relu_deriv)
39 delta4 = backwards(delta3, params, 'layer2', relu_deriv)
40 delta5 = backwards(delta4, params, 'layer1', relu_deriv)
41 # apply gradient
42 # for layers
43 for name in ('output','layer3','layer2', 'layer1'):
44     grad_W = params['grad_W' + name]
45     grad_b = params['grad_b' + name]
46     params['W' + name] -= learning_rate * grad_W
47     params['b' + name] -= learning_rate * grad_b
```

Q5.1.2 (Extra Credit) Code [5 points] To help even more with convergence speed, we will implement [momentum](#). Now, instead of updating $W = W - \alpha \frac{\partial J}{\partial W}$, we will use the update rules $M_W = 0.9M_W - \alpha \frac{\partial J}{\partial W}$ and $W = W + M_W$. To implement this, populate the parameters dictionary with zero-initialized momentum accumulators, one for each parameter. Then simply perform both update equations for every batch. **Include your code in the writeup.**

Q5.1.2


```
1 # initialize layers here
2 #####
3 ##### your code here #####
4 #####
5
6 initialize_weights(1024, 32, params, "layer1")
7 # zero-initialized momentum accumulators
8 params['Mw' + 'layer1'] = np.zeros((1024,32))
9 params['Mb' + 'layer1'] = np.zeros((32))
10
11 initialize_weights(32, 32, params, "layer2")
12 params['Mw' + 'layer2'] = np.zeros((32,32))
13 params['Mb' + 'layer2'] = np.zeros((32))
14
15 initialize_weights(32, 32, params, "layer3")
16 params['Mw' + 'layer3'] = np.zeros((32,32))
17 params['Mb' + 'layer3'] = np.zeros((32))
18
19 initialize_weights(32, 1024, params, "output")
20 params['Mw' + 'output'] = np.zeros((32,1024))
21 params['Mb' + 'output'] = np.zeros((1024))
22 # should look like your previous training loops
23
24 #####
25 ##### your code here #####
26 #####
27 ...
28 # apply gradient
29 # for layers
30 for name in ('output','layer3','layer2', 'layer1'):
31     grad_W = params['grad_W' + name]
32     grad_b = params['grad_b' + name]
33     params['Mw' + name] = 0.9 * params['Mw' + name] - learning_rate * grad_W
34     params['W' + name] = params['W' + name] + params['Mw' + name]
35     params['Mb' + name] = 0.9 * params['Mb' + name] - learning_rate * grad_b
36     params['b' + name] = params['b' + name] + params['Mb' + name]
```

5.2 Training the Autoencoder


Q5.2 (Extra Credit) Writeup/Code [5 points] Using the provided default settings, train the network for 100 epochs. Plot the training loss curve and include it in the writeup. What do you observe?

Q5.2

Not using momentum:



Use momentum:




The loss get decreased and become stable when meet one epoch value

5.3 Evaluating the Autoencoder


Q5.3.1 (Extra Credit) Writeup/Code [5 points] Now let's evaluate how well the autoencoder has been trained. Select 5 classes from the total 36 classes in the validation set and for each selected class include in your report 2 validation images and their reconstruction. What differences do you observe in the reconstructed validation images compared to the original ones?

Q5.3.1

Result: Not using momentum:



Use momentum:



Difference:

The reconstructed ones Look more blurry and the boundaries are not as clear as the original image

Code:

```
1 h = forward(visualize_x, params, 'layer1',relu)
2 h = forward(h, params, 'layer2',relu)
3 h = forward(h, params, 'layer3',relu)
4 reconstructed_x = forward(h, params, 'output', sigmoid)
```

Q5.3.2 (Extra Credit) Writeup [5 points] Let's evaluate the reconstruction quality using Peak Signal-to-noise Ratio (PSNR). PSNR is defined as

$$\text{PSNR} = 20 \times \log_{10}(\text{MAX}_I) - 10 \times \log_{10}(\text{MSE}) \quad (1)$$

where MAX_I is the maximum possible pixel value of the image, and MSE (mean squared error) is computed across all pixels. Said another way, maximum refers to the brightest overall sum (maximum positive value of the sum). You may use `skimage.metrics.peak_signal_noise_ratio` for convenience. Report the average PSNR you get from the autoencoder across all images in the validation set (it should be around 15).

Q5.3.2

Result:

Not using momentum:

Average PSNR : 13.379990823521801

using momentum:

Average PSNR: 11.322928512093558

```
1 h = forward(valid_x, params, 'layer1',relu)
2 h = forward(h, params, 'layer2',relu)
3 h = forward(h, params, 'layer3',relu)
4 result = forward(h, params, 'output', sigmoid)
5 rsnr_total=0
6 for i in range(len(valid_x)):
7     rsnr_total+=peak_signal_noise_ratio(result[i],valid_x[i])
8 print("Average PSNR: "+str(rsnr_total/len(valid_x)))
```

6 PyTorch

While you were able to derive manual backpropagation rules for sigmoid and fully-connected layers, wouldn't it be nice if someone did that for lots of useful primitives and made it

fast and easy to use for general computation? Meet [automatic differentiation](#). Since we have high-dimensional inputs (images) and low-dimensional outputs (a scalar loss), it turns out **forward mode AD** is very efficient. Popular autodiff packages include [pytorch](#) (Facebook), [tensorflow](#) (Google), [autograd](#) (Boston-area academics). Autograd provides its own replacement for numpy operators and is a drop-in replacement for numpy, except you can ask for gradients now. The other two are able to utilize GPUs to perform highly optimized and parallel computations, and are very popular for researchers who train large networks. Tensorflow asks you to build a computational graph using its API, and then is able to pass data through that graph. PyTorch builds a dynamic graph and allows you to mix autograd functions with normal python code much more smoothly, so it is currently more popular in academia.

We will use [PyTorch](#) as a framework. Many computer vision projects use neural networks as a basic building block, so familiarity with one of these frameworks is a good skill to develop. Here, we basically replicate and slightly expand our handwritten character recognition networks, but do it in PyTorch instead of doing it ourselves. Feel free to use any tutorial you like, but we like [the official one](#) or [this tutorial](#) (in a jupyter notebook) or [these slides](#) (starting from number 35).



For this section, you're free to implement these however you like. All of the tasks required here are fairly small and don't require a GPU if you use small networks. Include the neural network code for each part in the writeup.

6.1 Train a neural network in PyTorch

Q6.1.1 Code/Writeup [5 points] Re-write and re-train your **fully-connected** network on the included **NIST36** in PyTorch. Plot training accuracy and loss over time.

Q6.1.1

plot:acc(left) loss(right)



(please see code in next page)

runq6.1.1.py

```
1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4 import scipy.io
5 from torch.utils.data import Dataset,DataLoader
6 import matplotlib.pyplot as plt
7
8
9 def plot(train_loss, valid_loss, train_acc, valid_acc):
10     # plot loss curves
11     plt.plot(range(len(train_loss)), train_loss, label="training")
12     plt.plot(range(len(valid_loss)), valid_loss, label="validation")
13     plt.xlabel("epoch")
14     plt.ylabel("average loss")
15     plt.xlim(0, len(train_loss) - 1)
16     plt.ylim(0, None)
17     plt.legend()
18     plt.grid()
19     plt.show()
20
21     # plot accuracy curves
22     plt.plot(range(len(train_acc)), train_acc, label="training")
23     plt.plot(range(len(valid_acc)), valid_acc, label="validation")
24     plt.xlabel("epoch")
25     plt.ylabel("accuracy")
26     plt.xlim(0, len(train_acc) - 1)
27     plt.ylim(0, None)
28     plt.legend()
29     plt.grid()
30     plt.show()
31
32 class Dataset(Dataset):
33     def __init__(self, x, y):
34         self.x = x
35         self.y = y
36
37     def __getitem__(self, index):
38         x = torch.Tensor(self.x[index])
39         y = torch.Tensor(self.y[index])
40         return (x, y)
41
42     def __len__(self):
43         count = self.x.shape[0]
44         return count
```

```

45
46
47 class FC_Net(nn.Module):
48     def __init__(self, size_in, size_out):
49         super().__init__()
50         self.fc1 = nn.Linear(size_in,64)
51         self.fc2 = nn.Linear(64, size_out)
52
53     def forward(self,x):
54         x = self.fc1(x)
55         x = F.sigmoid(x)
56         out = self.fc2(x)
57         return out
58
59
60 train_data = scipy.io.loadmat('../data/nist36_train.mat')
61 valid_data = scipy.io.loadmat('../data/nist36_valid.mat')
62 test_data = scipy.io.loadmat('../data/nist36_test.mat')
63 train_x, train_y = train_data['train_data'], train_data['train_labels']
64 valid_x, valid_y = valid_data['valid_data'], valid_data['valid_labels']
65 test_x, test_y = test_data['test_data'], test_data['test_labels']
66
67 # model and loss, optimizer
68 model = FC_Net(train_x.shape[1], train_y.shape[1])
69 print(model)
70 loss_fn = nn.CrossEntropyLoss()
71 optimizer = torch.optim.Adam(model.parameters(), lr=0.002)
72
73 # load data, shuffle is necessary!!!
74 train_data = Dataset(train_x,train_y)
75 train_loader = DataLoader(train_data, batch_size=100, shuffle=True)
76
77 test_data = Dataset(test_x,test_y)
78 test_loader = DataLoader(test_data, batch_size=100)
79
80 valid_data = Dataset(valid_x,valid_y)
81 valid_loader = DataLoader(valid_data, batch_size=100)
82
83 epoch = 50
84 train_loss = []
85 valid_loss = []
86 train_acc = []
87 valid_acc = []
88 for i in range(epoch):
89     # training part
90     print("Training iteration: " + str(i+1))

```

```


91     size = len(train_loader.dataset)
92     num_batches = len(train_loader)
93     model.train()
94     loss_val, correct = 0, 0
95     for data in train_loader:
96         x, y = data
97         pred = model(x)
98         loss = loss_fn(pred, y)
99         #pred = nn.Softmax(dim=1)(pred)
100        correct += (pred.argmax(1) ==
101                     y.argmax(1)).type(torch.float).sum().item()
102        loss_val+=loss.item()
103
104        # Backpropagation
105        optimizer.zero_grad()
106        loss.backward()
107        optimizer.step()
108        loss_val /= num_batches
109        correct /= size
110        train_loss.append(loss_val)
111        train_acc.append(correct)
112        print("Train loss: " + str(loss_val) + ", Accuracy: " + str(correct *
113                                     100) + "%")
114
115        # validation part
116        size = len(valid_loader.dataset)
117        num_batches = len(valid_loader)
118        model.eval()
119        loss, correct = 0, 0
120        with torch.no_grad():
121            for X, y in valid_loader:
122                pred = model(X)
123                loss += loss_fn(pred, y).item()
124                #pred = nn.Softmax(dim=1)(pred)
125                correct += (pred.argmax(1) ==
126                             y.argmax(1)).type(torch.float).sum().item()
127                loss /= num_batches
128                correct /= size
129                valid_loss.append(loss)
130                valid_acc.append(correct)
131                print("Validation loss: "+str(loss)+", Accuracy: " +
132                      str(correct*100)+"%")
133
134    plot(train_loss, valid_loss, train_acc, valid_acc)

```

Q6.1.2 Code/Writeup [5 points] Train a **convolutional** neural network with PyTorch on the included **NIST36** dataset. Compare its performance with the previous fully-connected network.

Q6.1.2

plot:acc(left) loss(right)



From plot we can see: The result (accuracy around 90%) is much more better than previous fully-connected network (accuracy around 78%)

runq6.1.2.py

```
1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4 import scipy.io
5 from torch.utils.data import Dataset, DataLoader
6 import matplotlib.pyplot as plt
7
8
9 def plot(train_loss, valid_loss, train_acc, valid_acc):
10     # plot loss curves
11     plt.plot(range(len(train_loss)), train_loss, label="training")
12     plt.plot(range(len(valid_loss)), valid_loss, label="validation")
13     plt.xlabel("epoch")
14     plt.ylabel("average loss")
15     plt.xlim(0, len(train_loss) - 1)
16     plt.ylim(0, None)
17     plt.legend()
18     plt.grid()
19     plt.show()
```

```

20
21     # plot accuracy curves
22     plt.plot(range(len(train_acc)), train_acc, label="training")
23     plt.plot(range(len(valid_acc)), valid_acc, label="validation")
24     plt.xlabel("epoch")
25     plt.ylabel("accuracy")
26     plt.xlim(0, len(train_acc) - 1)
27     plt.ylim(0, None)
28     plt.legend()
29     plt.grid()
30     plt.show()

31
32
33 class Dataset(Dataset):
34     def __init__(self, x, y):
35         self.x = x
36         self.y = y
37
38     def __getitem__(self, index):
39         x = torch.Tensor(self.x[index])
40         y = torch.Tensor(self.y[index])
41         return (x, y)
42
43     def __len__(self):
44         count = self.x.shape[0]
45         return count
46
47
48 class CNN_Net(nn.Module):
49     def __init__(self, size_in, size_out):
50         super().__init__()
51         self.cnn_model = nn.Sequential(
52             nn.Conv2d(1, 6, 5),
53             nn.ReLU(),
54             nn.MaxPool2d(2,2),
55             nn.Conv2d(6, 16, 5),
56             nn.ReLU(),
57             nn.MaxPool2d(2,2),
58             nn.Flatten(),
59             nn.Linear(16 * 5 * 5, 128),
60             nn.ReLU(),
61             nn.Linear(128, 64),
62             nn.ReLU(),
63             nn.Linear(64, size_out)
64         )
65

```

```

66     def forward(self,x):
67         out = self.cnn_model(x)
68         return out
69
70
71 train_data = scipy.io.loadmat('../data/nist36_train.mat')
72 valid_data = scipy.io.loadmat('../data/nist36_valid.mat')
73 test_data = scipy.io.loadmat('../data/nist36_test.mat')
74 train_x, train_y = train_data['train_data'], train_data['train_labels']
75 valid_x, valid_y = valid_data['valid_data'], valid_data['valid_labels']
76 test_x, test_y = test_data['test_data'], test_data['test_labels']
77
78 # model and loss, optimizer
79 model = CNN_Net(train_x.shape[1], train_y.shape[1])
80 print(model)
81 loss_fn = nn.CrossEntropyLoss()
82 optimizer = torch.optim.Adam(model.parameters(), lr=0.002)
83
84 # load data, shuffle is necessary!!!
85 train_data = Dataset(train_x,train_y)
86 print(train_x.shape,train_y.shape)
87 train_loader = DataLoader(train_data, batch_size=100, shuffle=True)
88
89 test_data = Dataset(test_x,test_y)
90 test_loader = DataLoader(test_data, batch_size=100)
91
92 valid_data = Dataset(valid_x,valid_y)
93 valid_loader = DataLoader(valid_data, batch_size=100)
94
95 epoch = 50
96 train_loss = []
97 valid_loss = []
98 train_acc = []
99 valid_acc = []
100 for i in range(epoch):
101     # training part
102     print("Training iteration: " + str(i+1))
103     size = len(train_loader.dataset)
104     num_batches = len(train_loader)
105     model.train()
106     loss_val, correct = 0, 0
107     for data in train_loader:
108         x, y = data
109         N = x.shape[0]
110         # N,channel,h,w
111         x = x.reshape(N,1,32,32)

```

```

112     pred = model(x)
113     loss = loss_fn(pred, y)
114     #pred = nn.Softmax(dim=1)(pred)
115     correct += (pred.argmax(1) ==
116                   y.argmax(1)).type(torch.float).sum().item()
117     loss_val+=loss.item()

118     # Backpropagation
119     optimizer.zero_grad()
120     loss.backward()
121     optimizer.step()
122     loss_val /= num_batches
123     correct /= size
124     train_loss.append(loss_val)
125     train_acc.append(correct)
126     print("Train loss: " + str(loss_val) + ", Accuracy: " + str(correct * 
127           100) + "%")

128
129     # validation part
130     size = len(valid_loader.dataset)
131     num_batches = len(valid_loader)
132     model.eval()
133     loss, correct = 0, 0
134     with torch.no_grad():
135         for X, y in valid_loader:
136             N = X.shape[0]
137             # N,channel,h,w
138             X = X.reshape(N, 1, 32, 32)
139             pred = model(X)
140             loss += loss_fn(pred, y).item()
141             #pred = nn.Softmax(dim=1)(pred)
142             correct += (pred.argmax(1) ==
143                           y.argmax(1)).type(torch.float).sum().item()
144             loss /= num_batches
145             correct /= size
146             valid_loss.append(loss)
147             valid_acc.append(correct)
148             print("Validation loss: "+str(loss)+", Accuracy: " +
149               str(correct*100)+"%")

150     plot(train_loss, valid_loss, train_acc, valid_acc)



```

Q6.1.3 Code/Writeup [5 points] Train a **convolutional** neural network with PyTorch on **CIFAR-10** (`torchvision.datasets.CIFAR10`). Plot training accuracy and loss over

time.

Q6.1.3

plot:acc(left) loss(right)



runq6.1.3.py

```
1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4 import scipy.io
5 from torch.utils.data import Dataset,DataLoader
6 import matplotlib.pyplot as plt
7 import torchvision
8
9
10 def plot(train_loss, valid_loss, train_acc, valid_acc):
11     # plot loss curves
12     plt.plot(range(len(train_loss)), train_loss, label="training")
13     plt.plot(range(len(valid_loss)), valid_loss, label="validation")
14     plt.xlabel("epoch")
15     plt.ylabel("average loss")
16     plt.xlim(0, len(train_loss) - 1)
17     plt.ylim(0, None)
18     plt.legend()
19     plt.grid()
20     plt.show()
21
```

```

22     # plot accuracy curves
23     plt.plot(range(len(train_acc)), train_acc, label="training")
24     plt.plot(range(len(valid_acc)), valid_acc, label="validation")
25     plt.xlabel("epoch")
26     plt.ylabel("accuracy")
27     plt.xlim(0, len(train_acc) - 1)
28     plt.ylim(0, None)
29     plt.legend()
30     plt.grid()
31     plt.show()

32

33

34 class Dataset(Dataset):
35     def __init__(self, x, y):
36         self.x = x
37         self.y = y
38
39     def __getitem__(self, index):
40         x = torch.Tensor(self.x[index])
41         y = torch.Tensor(self.y[index])
42         return (x, y)
43
44     def __len__(self):
45         count = self.x.shape[0]
46         return count
47
48

49 class CNN_Net(nn.Module):
50     def __init__(self):
51         super().__init__()
52         self.cnn_model = nn.Sequential(
53             nn.Conv2d(3, 6, 5),
54             nn.ReLU(),
55             nn.MaxPool2d(2,2),
56             nn.Conv2d(6, 16, 5),
57             nn.ReLU(),
58             nn.MaxPool2d(2,2),
59             nn.Flatten(),
60             nn.Linear(16 * 5 * 5, 128),
61             nn.ReLU(),
62             nn.Linear(128, 64),
63             nn.ReLU(),
64             nn.Linear(64, 10)
65         )
66
67     def forward(self,x):

```

```

68         out = self.cnn_model(x)
69         return out
70
71
72 train_data = torchvision.datasets.CIFAR10(root='./data', train=True,
73                                         download=True, transform=torchvision.transforms.ToTensor())
73 test_data = torchvision.datasets.CIFAR10(root='./data', train=False,
74                                         download=True, transform=torchvision.transforms.ToTensor())
75
76 # model and loss, optimizer
77 model = CNN_Net()
78 print(model)
79 loss_fn = nn.CrossEntropyLoss()
80 optimizer = torch.optim.Adam(model.parameters(), lr=0.002)
81
82 # load data, shuffle is necessary!!!
82 train_loader = DataLoader(train_data, batch_size=64, shuffle=True)
83 valid_loader = DataLoader(test_data, batch_size=64)
84
85
86 epoch = 50
87 train_loss = []
88 valid_loss = []
89 train_acc = []
90 valid_acc = []
91 for i in range(epoch):
92     # training part
93     print("Training iteration: " + str(i+1))
94     size = len(train_loader.dataset)
95     num_batches = len(train_loader)
96     model.train()
97     loss_val, correct = 0, 0
98     for data in train_loader:
99         x, y = data
100        pred = model(x)
101        loss = loss_fn(pred, y)
102        #pred = nn.Softmax(dim=1)(pred)
103        correct += (pred.argmax(1) == y).type(torch.float).sum().item()
104        loss_val+=loss.item()
105
106        # Backpropagation
107        optimizer.zero_grad()
108        loss.backward()
109        optimizer.step()
110        loss_val /= num_batches
111        correct /= size

```

```



112     train_loss.append(loss_val)
113     train_acc.append(correct)
114     print("Train loss: " + str(loss_val) + ", Accuracy: " + str(correct *
115         ↪ 100) + "%")
116
117     # validation part
118     size = len(valid_loader.dataset)
119     num_batches = len(valid_loader)
120     model.eval()
121     loss, correct = 0, 0
122     with torch.no_grad():
123         for X, y in valid_loader:
124             pred = model(X)
125             loss += loss_fn(pred, y).item()
126             #pred = nn.Softmax(dim=1)(pred)
127             correct += (pred.argmax(1) ==
128                         ↪ y).type(torch.float).sum().item()
129             loss /= num_batches
130             correct /= size
131             valid_loss.append(loss)
132             valid_acc.append(correct)
133             print("Validation loss: "+str(loss)+", Accuracy: " +
134                 ↪ str(correct*100)+"%")
135
136     plot(train_loss, valid_loss, train_acc, valid_acc)

```

Q6.1.4 Code/Writeup [10 points] In Homework 1, we tried scene classification with the bag-of-words (BoW) approach on a subset of the **SUN database**. Use the same dataset in HW1, and implement a **convolutional** neural network with PyTorch for **scene** classification. Compare your result with the one you got in HW1, and briefly comment on it.

Q6.1.4

plot:acc(left) loss(right)



We can see that the accuracy is around 50%, it's not as good as HW1 method. So for scene classification, only use basic CNN is not enough. Some other methods like data transforming are still needed. OR better CNN can be built

runq6.1.4.py

```
1 import os.path
2
3 import skimage
4 import torch
5 import torch.nn as nn
6 from PIL import Image
7 from torch.utils.data import Dataset,DataLoader
8 import matplotlib.pyplot as plt
9 import numpy as np
10
11
12
13 def plot(train_loss, valid_loss, train_acc, valid_acc):
14     # plot loss curves
15     plt.plot(range(len(train_loss)), train_loss, label="training")
16     plt.plot(range(len(valid_loss)), valid_loss, label="validation")
17     plt.xlabel("epoch")
18     plt.ylabel("average loss")
19     plt.xlim(0, len(train_loss) - 1)
20     plt.ylim(0, None)
21     plt.legend()
22     plt.grid()
```

```

23     plt.show()

24
25     # plot accuracy curves
26     plt.plot(range(len(train_acc)), train_acc, label="training")
27     plt.plot(range(len(valid_acc)), valid_acc, label="validation")
28     plt.xlabel("epoch")
29     plt.ylabel("accuracy")
30     plt.xlim(0, len(train_acc) - 1)
31     plt.ylim(0, None)
32     plt.legend()
33     plt.grid()
34     plt.show()

35
36 img_size = 64

37
38 class Dataset(Dataset):
39     def __init__(self, x, y):
40         self.x = x
41         self.y = y
42
43     def __getitem__(self, index):
44         x = torch.Tensor(self.x[index])
45         y = torch.Tensor(self.y[index])
46         return (x, y)
47
48     def __len__(self):
49         count = self.x.shape[0]
50         return count

51
52
53 class CNN_Net(nn.Module):
54     def __init__(self):
55         super().__init__()
56         self.cnn_model = nn.Sequential(
57             nn.Conv2d(1, 16, 5),
58             nn.ReLU(),
59             nn.MaxPool2d(2,2),
60
61             nn.Conv2d(16, 32, 5),
62             nn.ReLU(),
63             nn.MaxPool2d(2,2),
64
65             nn.Conv2d(32, 64, 5),
66             nn.ReLU(),
67             nn.MaxPool2d(2, 2),
68

```

```

69         nn.Flatten(),
70         nn.Linear(1024, 128),
71         nn.ReLU(),
72         nn.Linear(128, 64),
73         nn.ReLU(),
74         nn.Linear(64, 8)
75     )
76
77     def forward(self,x):
78         out = self.cnn_model(x)
79         return out
80
81
82     def get_data(path,x_txt,y_txt):
83         # for img
84         f = open(x_txt, 'r')
85         lines = f.readlines()
86         #print(lines)
87         f.close()
88         x = []
89         for line in lines:
90             line = line.replace("\n", " ")
91             img_path = path+line
92             img = Image.open(img_path)
93             img = img.resize((img_size, img_size))
94             img = np.array(img.convert('L'))
95             img=img/255.0
96             x.append(img)
97
98         f = open(y_txt, 'r')
99         lines = f.readlines()
100        f.close()
101        y = []
102        for line in lines:
103            line = line.replace("\n", " ")
104            y.append(int(line))
105
106        y = np.array(y)
107        # print(y)
108        x = np.array(x)
109        # N = x.shape[0]
110        # # N,channel,h,w
111        # x = x.reshape(N, 1, 64, 64)
112        y=y.reshape(-1,1)
113        # print(y)
114        print(x.shape)

```

```

115     print(y.shape)
116     return x,y
117
118
119 # read
120 train_x_path = '../data/SUN/train_files.txt'
121 train_y_path = '../data/SUN/train_labels.txt'
122 test_x_path = '../data/SUN/test_files.txt'
123 test_y_path = '../data/SUN/test_labels.txt'
124 path = '../data/SUN/'
125 # save
126 x_train_save = '../data/sun_train_x.npy'
127 y_train_save = '../data/sun_train_y.npy'
128 x_test_save = '../data/sun_test_x.npy'
129 y_test_save = '../data/sun_test_y.npy'
130
131 if os.path.exists(x_train_save):
132     train_x = np.load(x_train_save)
133     train_y = np.load(y_train_save)
134     valid_x = np.load(x_test_save)
135     valid_y = np.load(y_test_save)
136 else:
137     train_x, train_y = get_data(path, train_x_path, train_y_path)
138     valid_x, valid_y = get_data(path, test_x_path, test_y_path)
139     np.save(x_train_save,train_x)
140     np.save(y_train_save, train_y)
141     np.save(x_test_save, valid_x)
142     np.save(y_test_save, valid_y)
143
144 # model and loss, optimizer
145 model = CNN_Net()
146 print(model)
147 loss_fn = nn.CrossEntropyLoss()
148 optimizer = torch.optim.Adam(model.parameters(), lr=0.002)
149
150 # load data, shuffle is necessary!!!
151 train_data = Dataset(train_x,train_y)
152 train_loader = DataLoader(train_data, batch_size=32, shuffle=True)
153
154 valid_data = Dataset(valid_x,valid_y)
155 valid_loader = DataLoader(valid_data, batch_size=32)
156
157 epoch = 30
158 train_loss = []
159 valid_loss = []
160 train_acc = []

```

```

161 valid_acc = []
162
163 for i in range(epoch):
164     # training part
165     print("Training iteration: " + str(i+1))
166     size = len(train_loader.dataset)
167     num_batches = len(train_loader)
168     model.train()
169     loss_val, correct = 0, 0
170     for data in train_loader:
171         x, y = data
172         N = x.shape[0]
173         # N,channel,h,w
174         x = x.reshape(N,1,img_size,img_size)
175         y=y.reshape(-1)
176         #print(y.dtype)
177         pred = model(x)
178         y = y.long()
179
180         loss = loss_fn(pred, y)
181         #pred = nn.Softmax(dim=1)(pred)
182         correct += (pred.argmax(1) == y).type(torch.float).sum().item()
183         loss_val+=loss.item()
184
185         # Backpropagation
186         optimizer.zero_grad()
187         loss.backward()
188         optimizer.step()
189         loss_val /= num_batches
190         correct /= size
191         train_loss.append(loss_val)
192         train_acc.append(correct)
193         print("Train loss: " + str(loss_val) + ", Accuracy: " + str(correct *
194             100) + "%")
195
196     # validation part
197     size = len(valid_loader.dataset)
198     num_batches = len(valid_loader)
199     model.eval()
200     loss, correct = 0, 0
201     with torch.no_grad():
202         for X, y in valid_loader:
203             N = X.shape[0]
204             # N,channel,h,w
205             X = X.reshape(N, 1, img_size, img_size)

```

```

206     pred = model(X)
207     y = y.reshape(-1)
208     y = y.long()
209     loss += loss_fn(pred, y).item()
210     #pred = nn.Softmax(dim=1)(pred)
211     correct += (pred.argmax(1) ==
212                   y).type(torch.float).sum().item()
213     loss /= num_batches
214     correct /= size
215     valid_loss.append(loss)
216     valid_acc.append(correct)
217     print("Validation loss: "+str(loss)+", Accuracy: " +
218           str(correct*100)+"%")
219
220 plot(train_loss, valid_loss, train_acc, valid_acc)

```

6.2 Fine Tuning

When training from scratch, a lot of epochs and data are often needed to learn anything meaningful. One way to avoid this is to instead initialize the weights more intelligently.



These days, it is most common to initialize a network with weights from another deep network that was trained for a different purpose. This is because, whether we are doing image classification, segmentation, recognition etc..., most real images share common properties. Simply copying the weights from the other network to yours gives your network a head start, so your network does not need to learn these common weights from scratch all over again. This is also referred to as fine tuning.

Q6.2 Code/Writeup [5 points] Fine-tune a single layer classifier using pytorch on the [flowers 17](#) (or [flowers 102!](#)) dataset using [squeezenet1_1](#), as well as an architecture you've designed yourself (for example 3 convolutional layers followed 2 fully connected layers, it's standard [slide 6](#)) and trained from scratch. How do they compare?



We include a script in `scripts/` to fetch the flowers dataset and extract it in a way that `torchvision.datasets.ImageFolder` can consume it, see [an example](#), from `data/oxford-flowers17`. You should look at how SqueezeNet is [defined](#), and just replace the classifier layer. There exists a pretty good example for [fine-tuning](#) in PyTorch.

Q6.2

plot: My network



plot: Use fine tuning of queezenet1_1 can be much better.



We can see that using my network is not as good as fine tuning and even cause overfitting when epoch is large.

So using fine tuning

6.2.2.py

```
1 import torch
2 import torchvision
3 from torchvision import transforms, datasets
4 import os.path
5
6 import skimage
7 import torch
8 import torch.nn as nn
9 from PIL import Image
10 from torch.utils.data import Dataset, DataLoader
```

```

11 import matplotlib.pyplot as plt
12 import numpy as np
13 img_size = 32
14 input_size = 224
15 def plot(train_loss, valid_loss, train_acc, valid_acc):
16     # plot loss curves
17     plt.plot(range(len(train_loss)), train_loss, label="training")
18     plt.plot(range(len(valid_loss)), valid_loss, label="validation")
19     plt.xlabel("epoch")
20     plt.ylabel("average loss")
21     plt.xlim(0, len(train_loss) - 1)
22     plt.ylim(0, None)
23     plt.legend()
24     plt.grid()
25     plt.show()
26
27     # plot accuracy curves
28     plt.plot(range(len(train_acc)), train_acc, label="training")
29     plt.plot(range(len(valid_acc)), valid_acc, label="validation")
30     plt.xlabel("epoch")
31     plt.ylabel("accuracy")
32     plt.xlim(0, len(train_acc) - 1)
33     plt.ylim(0, None)
34     plt.legend()
35     plt.grid()
36     plt.show()
37
38 data_transforms = transforms.Compose([
39     transforms.RandomResizedCrop(input_size),
40     transforms.RandomHorizontalFlip(),
41     transforms.ToTensor(),
42     transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
43 ])
44 trans_dataset =
45     datasets.ImageFolder(root='../data/oxford-flowers17/train',
46                         transform=data_transforms)
47 train_loader = torch.utils.data.DataLoader(trans_dataset,
48                                             batch_size=4, shuffle=True)
49 trans_dataset = datasets.ImageFolder(root='../data/oxford-flowers17/val',
50                                     transform=data_transforms)
51 valid_loader = torch.utils.data.DataLoader(trans_dataset,
52                                         batch_size=4, shuffle=True)
53
54     # model and loss, optimizer
55     model = torchvision.models.squeezenet1_1(pretrained=True)
56     feature_extract = True

```

```

56 # set false
57 for param in model.parameters():
58     param.requires_grad = False
59 # change classifier layer
60 model.classifier[1] = nn.Conv2d(512, 17, kernel_size=(1,1), stride=(1,1))
61 model.num_classes = 17
62 input_size = 224
63 print(model)
64 loss_fn = nn.CrossEntropyLoss()
65 # select parameters
66 params_to_update = model.parameters()
67 print("Params to learn:")
68 # only change
69 if feature_extract:
70     params_to_update = []
71     for name,param in model.named_parameters():
72         if param.requires_grad == True:
73             params_to_update.append(param)
74             print("\t",name)
75 else:
76     for name,param in model.named_parameters():
77         if param.requires_grad == True:
78             print("\t",name)
79
80 # Observe that all parameters are being optimized
81 optimizer = torch.optim.SGD(params_to_update, lr=0.001)
82
83 epoch = 30
84 train_loss = []
85 valid_loss = []
86 train_acc = []
87 valid_acc = []
88
89 for i in range(epoch):
90     # training part
91     print("Training iteration: " + str(i+1))
92     size = len(train_loader.dataset)
93     num_batches = len(train_loader)
94     model.train()
95     loss_val, correct = 0, 0
96     for data in train_loader:
97         x, y = data
98         pred = model(x)
99         loss = loss_fn(pred, y)
100        #pred = nn.Softmax(dim=1)(pred)
101        correct += (pred.argmax(1) == y).type(torch.float).sum().item()

```

```

102         loss_val+=loss.item()
103
104     # Backpropagation
105     optimizer.zero_grad()
106     loss.backward()
107     optimizer.step()
108     loss_val /= num_batches
109     correct /= size
110     train_loss.append(loss_val)
111     train_acc.append(correct)
112     print("Train loss: " + str(loss_val) + ", Accuracy: " + str(correct *
113                               ↵ 100) + "%")
114
115     # validation part
116     size = len(valid_loader.dataset)
117     num_batches = len(valid_loader)
118     model.eval()
119     loss, correct = 0, 0
120     with torch.no_grad():
121         for X, y in valid_loader:
122             pred = model(X)
123             loss += loss_fn(pred, y).item()
124             #pred = nn.Softmax(dim=1)(pred)
125             correct += (pred.argmax(1) ==
126                         ↵ y).type(torch.float).sum().item()
127             loss /= num_batches
128             correct /= size
129             valid_loss.append(loss)
130             valid_acc.append(correct)
131             print("Validation loss: "+str(loss)+", Accuracy: " +
132                               ↵ str(correct*100)+"%")
133
134     plot(train_loss, valid_loss, train_acc, valid_acc)

```

6.2.1.py

```

1 import torch
2 from torchvision import transforms, datasets
3 import os.path
4
5 import skimage
6 import torch
7 import torch.nn as nn
8 from PIL import Image

```

```

9  from torch.utils.data import Dataset,DataLoader
10 import matplotlib.pyplot as plt
11 import numpy as np
12 img_size = 64
13
14 def plot(train_loss, valid_loss, train_acc, valid_acc):
15     # plot loss curves
16     plt.plot(range(len(train_loss)), train_loss, label="training")
17     plt.plot(range(len(valid_loss)), valid_loss, label="validation")
18     plt.xlabel("epoch")
19     plt.ylabel("average loss")
20     plt.xlim(0, len(train_loss) - 1)
21     plt.ylim(0, None)
22     plt.legend()
23     plt.grid()
24     plt.show()
25
26     # plot accuracy curves
27     plt.plot(range(len(train_acc)), train_acc, label="training")
28     plt.plot(range(len(valid_acc)), valid_acc, label="validation")
29     plt.xlabel("epoch")
30     plt.ylabel("accuracy")
31     plt.xlim(0, len(train_acc) - 1)
32     plt.ylim(0, None)
33     plt.legend()
34     plt.grid()
35     plt.show()
36
37 class CNN_Net(nn.Module):
38     def __init__(self):
39         super().__init__()
40         self.cnn_model = nn.Sequential(
41             nn.Conv2d(3, 16, 5),
42             nn.ReLU(),
43             nn.MaxPool2d(2,2),
44
45             nn.Conv2d(16, 32, 5),
46             nn.ReLU(),
47             nn.MaxPool2d(2,2),
48
49             nn.Conv2d(32, 64, 5),
50             nn.ReLU(),
51             nn.MaxPool2d(2, 2),
52
53             nn.Flatten(),
54             nn.Linear(1024, 128),

```

```

55         nn.ReLU(),
56         nn.Linear(128, 64),
57         nn.ReLU(),
58         nn.Linear(64, 17)
59     )
60
61     def forward(self,x):
62         out = self.cnn_model(x)
63         return out
64
65     data_transforms = transforms.Compose([
66         transforms.RandomHorizontalFlip(),
67         transforms.Resize((img_size, img_size)),
68         transforms.ToTensor(),
69         transforms.Normalize(mean=[0.485, 0.456, 0.406],
70                             std=[0.229, 0.224, 0.225])
71     ])
72     trans_dataset =
73     ↵ datasets.ImageFolder(root='./data/oxford-flowers17/train',
74                           transform=data_transforms)
75     train_loader = torch.utils.data.DataLoader(trans_dataset,
76                                               batch_size=4, shuffle=True)
76     trans_dataset = datasets.ImageFolder(root='./data/oxford-flowers17/val',
77                                         transform=data_transforms)
78     valid_loader = torch.utils.data.DataLoader(trans_dataset,
79                                               batch_size=4, shuffle=True)
80
81     # model and loss, optimizer
82     model = CNN_Net()
83     print(model)
84     loss_fn = nn.CrossEntropyLoss()
85     optimizer = torch.optim.Adam(model.parameters(), lr=0.002)
86
87     epoch = 30
88     train_loss = []
89     valid_loss = []
90     train_acc = []
91     valid_acc = []
92
93     for i in range(epoch):
94         # training part
95         print("Training iteration: " + str(i+1))
96         size = len(train_loader.dataset)
97         num_batches = len(train_loader)
98         model.train()
99         loss_val, correct = 0, 0

```

```

100     for data in train_loader:
101         x, y = data
102         pred = model(x)
103         loss = loss_fn(pred, y)
104         #pred = nn.Softmax(dim=1)(pred)
105         correct += (pred.argmax(1) == y).type(torch.float).sum().item()
106         loss_val+=loss.item()
107
108         # Backpropagation
109         optimizer.zero_grad()
110         loss.backward()
111         optimizer.step()
112         loss_val /= num_batches
113         correct /= size
114         train_loss.append(loss_val)
115         train_acc.append(correct)
116         print("Train loss: " + str(loss_val) + ", Accuracy: " + str(correct *
117             100) + "%")
118
119         # validation part
120         size = len(valid_loader.dataset)
121         num_batches = len(valid_loader)
122         model.eval()
123         loss, correct = 0, 0
124         with torch.no_grad():
125             for X, y in valid_loader:
126                 pred = model(X)
127                 loss += loss_fn(pred, y).item()
128                 #pred = nn.Softmax(dim=1)(pred)
129                 correct += (pred.argmax(1) ==
130                     y).type(torch.float).sum().item()
131                 loss /= num_batches
132                 correct /= size
133                 valid_loss.append(loss)
134                 valid_acc.append(correct)
135                 print("Validation loss: "+str(loss)+", Accuracy: " +
136                     str(correct*100)+"%")
137
138 plot(train_loss, valid_loss, train_acc, valid_acc)

```


6.3 Neural Networks in the Real World

Often, we train neural networks on standard datasets and evaluate on held out validation data. How would a model trained on ImageNet perform in the wild?

Q6.3 (Extra Credit) Neural Networks in the Real World [20 points] Download an ImageNet pretrained image classification model of your choice from [torchvision](#). Using this model, pick a single category (out of the 1000 used to train the model) and evaluate the validation performance of this category. You can download the ImageNet validation data from the [challenge page](#) by creating an account (top right). Torchvision has a [dataloader](#) to help you load and process ImageNet data automatically. Next, find an instance of this selected category in the real world and take a dynamic (i.e with some movement) video of this object. Extract all of the frames from this video and apply your pretrained model to each frame and compare the accuracy of the classifier on your video compared with the images in the validation set. Why might this be? Can you suggest ways to make your model more robust?

Q6.3

Create account and download validation data(test with 4 categories):



Test with categories:

Accuracy for category: 0 is 92.0


Accuracy for category: 1 is 92.0

Accuracy for category: 2 is 90.0

Also, use laptop(681) in dataset to test:

Accuracy for category: 681 is 54.0

Test with my video frames (also choose laptop 681):



Accuracy for category: 681 is 48.181818181818% (106.0/220)

We can see that the result in validation dataset is much better than video I take. Because the image size, scale and processing may be different and the model is trained to fit better for validation data.

Ways: Train model with more datasets including my own data, do things like data transforming to make it better.

(code in next page:)

6.3.py

```
1 import torch
2 import torchvision
3 from torchvision import transforms, datasets
4 import os.path
5
6 import skimage
7 import torch
8 import torch.nn as nn
9 from PIL import Image
10 from torch.utils.data import Dataset,DataLoader
11 import matplotlib.pyplot as plt
12 import numpy as np
13 from torchvision.models import resnet152,ResNet152_Weights
14 input_size = 224
15
16 data_transforms = transforms.Compose([
17     transforms.RandomResizedCrop(input_size),
18     transforms.RandomHorizontalFlip(),
19     transforms.ToTensor(),
20     transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
21 ])
22
23 val_data=torchvision.datasets.ImageNet(root="D:\CV", split='val',
24     transform=data_transforms)
25 val_data = DataLoader(val_data, batch_size=10)
26
27 # model and loss, optimizer
28 model = resnet152(weights=ResNet152_Weights.IMAGENET1K_V2)
29 model.eval()
30 k = 0
31 acc=0
32 sum=0
33 # tench and goldfish
34 with torch.no_grad():
35     for X, y in val_data:
36         if int(y[0])!=k:
37             # accuracy
38             res = (acc/sum)*100
39             print("Accuracy for category: "+str(k) +" is "+str(res)+"%")
40             # next category
41             k+=1
42             sum=0
43             acc=0
44     # only get accuracy for first 2
```

```

44     if k==3:
45         break
46     pred = model(X)
47     acc+=(pred.argmax(1) == y).type(torch.float).sum().item()
48     sum+=10

```

6.3real.py

```

1 import torch
2 import torchvision
3 from torchvision import transforms, datasets
4 import os.path
5
6 import skimage
7 import torch
8 import torch.nn as nn
9 from PIL import Image
10 from torch.utils.data import Dataset,DataLoader
11 import matplotlib.pyplot as plt
12 import numpy as np
13 from torchvision.models import resnet152,ResNet152_Weights
14 img_size = 240
15 input_size = 224
16
17 data_transforms = transforms.Compose([
18     transforms.Resize((img_size, img_size)),
19     transforms.CenterCrop((input_size, input_size)),
20     transforms.ToTensor(),
21     transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
22 ])
23
24 # car
25 real = 681
26 hymenoptera_dataset = datasets.ImageFolder(root='./data/real_data',
27                                              transform=data_transforms)
28 val_loader = torch.utils.data.DataLoader(hymenoptera_dataset,
29                                         batch_size=10, shuffle=True)
30
31 # model and loss, optimizer
32 model = resnet152(weights=ResNet152_Weights.IMAGENET1K_V2)
33 model.eval()
34 acc=0
35 sum=0
36 with torch.no_grad():
37     for X, y in val_loader:

```

```
38     pred = model(X)
39     acc+=(pred.argmax(1) == real).type(torch.float).sum().item()
40     sum+=10
41     res = (acc / sum) * 100
42     print("Accuracy for category: " + str(real) + " is " + str(res) +
43           "%")
```

6.3save.py

```
1 import cv2
2
3 cap = cv2.VideoCapture('../data/real.mp4')
4
5 i = 0
6 while (cap.isOpened()):
7     ret, frame = cap.read()
8     if ret == True:
9
10         filename = '../data/real_data/1/' + str(i) + '.jpg'
11         cv2.imwrite(filename, frame)
12         i+=1
13     else:
14         break
15
16 cap.release()
17 cv2.destroyAllWindows()
```

Deliverables

The assignment should be submitted to Gradescope. The writeup should be submitted as a pdf named <AndrewId>.pdf. The code should be submitted as a zip named <AndrewId>.zip. The zip when uncompressed should produce the following files.

- `nn.py`
- `q4.py`
- `run_q2.py`
- `run_q3.py`
- `run_q4.py`
- `run_q5.py` (extra credit)
- any `.py` files for Q6
- `util.py`

References

- [1] Xavier Glorot and Yoshua Bengio. Understanding the difficulty of training deep feedforward neural networks. 2010. <http://proceedings.mlr.press/v9/glorot10a/glorot10a.pdf>.
- [2] P. J. Grother. Nist special database 19 – handprinted forms and characters database. <https://www.nist.gov/srd/nist-special-database-19>, 1995.

Appendix: Neural Network Overview

Deep learning has quickly become one of the most applied machine learning techniques in computer vision. Convolutional neural networks have been applied to many different computer vision problems such as image classification, recognition, and segmentation with great success. In this assignment, you will first implement a fully connected feed-forward neural network for handwritten character classification. Then in the second part, you will implement a system to locate characters in an image, which you can then classify with your deep network. The end result will be a system that, given an image of handwritten text, will output the text contained in the image.

Mathematical overview

Here we will give a brief overview of the math for a single hidden layer feed-forward network. For a more detailed look at the math and derivation, please see the class slides.

A fully-connected network \mathbf{f} , for classification, applies a series of linear and non-linear functions to an input data vector \mathbf{x} of size $N \times 1$ to produce an output vector $\mathbf{f}(\mathbf{x})$ of size $C \times 1$, where each element i of the output vector represents the probability of \mathbf{x} belonging to the class i . Since the data samples are of dimensionality N , this means the input layer has N input units. To compute the value of the output units, we must first compute the values of all the hidden layers. The first hidden layer *pre-activation* $\mathbf{a}^{(1)}(\mathbf{x})$ is given by

$$\mathbf{a}^{(1)}(\mathbf{x}) = \mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)}$$

Then the *post-activation* values of the first hidden layer $\mathbf{h}^{(1)}(\mathbf{x})$ are computed by applying a non-linear activation function \mathbf{g} to the *pre-activation* values

$$\mathbf{h}^{(1)}(\mathbf{x}) = \mathbf{g}(\mathbf{a}^{(1)}(\mathbf{x})) = \mathbf{g}(\mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)})$$

Subsequent hidden layer ($1 < t \leq T$) pre- and post activations are given by:

$$\begin{aligned}\mathbf{a}^{(t)}(\mathbf{x}) &= \mathbf{W}^{(t)}\mathbf{h}^{(t-1)} + \mathbf{b}^{(t)} \\ \mathbf{h}^{(t)}(\mathbf{x}) &= \mathbf{g}(\mathbf{a}^{(t)}(\mathbf{x}))\end{aligned}$$

The output layer *pre-activations* $\mathbf{a}^{(T)}(\mathbf{x})$ are computed in a similar way

$$\mathbf{a}^{(T)}(\mathbf{x}) = \mathbf{W}^{(T)}\mathbf{h}^{(T-1)}(\mathbf{x}) + \mathbf{b}^{(T)}$$

and finally the *post-activation* values of the output layer are computed with

$$\mathbf{f}(\mathbf{x}) = \mathbf{o}(\mathbf{a}^{(T)}(\mathbf{x})) = \mathbf{o}(\mathbf{W}^{(T)}\mathbf{h}^{(T-1)}(\mathbf{x}) + \mathbf{b}^{(T)})$$

where \mathbf{o} is the output activation function. Please note the difference between \mathbf{g} and \mathbf{o} ! For this assignment, we will be using the sigmoid activation function for the hidden layer, so:

$$\mathbf{g}(y) = \frac{1}{1 + \exp(-y)}$$



Figure 2: Samples from NIST Special 19 dataset [2]

where when \mathbf{g} is applied to a vector, it is applied element wise across the vector.

Since we are using this deep network for classification, a common output activation function to use is the softmax function. This will allow us to turn the real value, possibly negative values of $\mathbf{a}^{(T)}(\mathbf{x})$ into a set of probabilities (vector of positive numbers that sum to 1). Letting \mathbf{x}_i denote the i^{th} element of the vector \mathbf{x} , the softmax function is defined as:

$$\mathbf{o}_i(\mathbf{y}) = \frac{\exp(\mathbf{y}_i)}{\sum_j \exp(\mathbf{y}_j)}$$

Gradient descent is an iterative optimisation algorithm, used to find the local optima. To find the local minima, we start at a point on the function and move in the direction of negative gradient (steepest descent) till some stopping criteria is met.

Backprop

The update equation for a general weight $W_{ij}^{(t)}$ and bias $b_i^{(t)}$ is

$$W_{ij}^{(t)} = W_{ij}^{(t)} - \alpha * \frac{\partial L_f}{\partial W_{ij}^{(t)}}(\mathbf{x}) \quad b_i^{(t)} = b_i^{(t)} - \alpha * \frac{\partial L_f}{\partial b_i^{(t)}}(\mathbf{x})$$

α is the learning rate. Please refer to the backpropagation slides for more details on how to derive the gradients. Note that here we are using softmax loss (which is different from the least square loss in the slides).

Debugging Checklist for Training Neural Networks

- Input-output pairs should make sense. Inspect them!
- Data loads correctly. Visualize it!
- Data is transformed correctly.
- Model and Data input-target dimensions must match
- Know what the output should look like. Is the result reasonable?
- Batch size is > 1 , else the gradient will be too noisy.
- Learning rate is tuned – not too small, not too large.
- Save the best checkpoint. Check on the validation set for overfitting!
- Run inference is in eval mode.
- Log everything!
- Fix the random seed in PyTorch / NumPy / OS when debugging
- Complex Error Logs are often Simple Bugs. Localize your errors in Tracebacks!
- Establish the invariants. What must be true? (Assert Statement)
- Make small changes between experiments to localize errors.
- Search Google/stack overflow with the error messages.